

Visualizing University of British Columbia Staff Salaries Split by Gender

Jade Bouchard

Table of contents

Aim	1
Data	1
Ethics	2
Methods	2
Gender Prediction	3
Exploratory Data Analysis	6
Limitations	11
Conclusion	12
References	13

Aim

This report explores The University of British Columbia (UBC) staff salaries based on guessed gender. When new salary data is released, re-runnnng the analysis will automatically update the figures and tables in this report with the latest data.

Data

Salary data for UBC staff members making over 75000 CAD annually was sourced from The University of British Columbia (2020). To access individual financial reports, click on the yearly links under the header [Statement of Financial Information \(SOFI\)](#).

Gender data was inferred using first names of staff members. In order to guess gender, I used baby name datasets (Statistics Canada, n.d.; Kaggle 2017; Sharma 2020).

Ethics

In this project first names are used to guess whether someone is “male” or “female”. Many people do not fit into these binary categories. In addition, while first names can sometimes be an indication of someone’s gender, first names are not inherently gendered. Misgendering in this project can have harmful effects.

I encourage anyone who notices a misgendering within this project to raise an issue in the issues tab on [my GitHub page](#), and it will be corrected. In addition, I encourage respectful and inclusive language in all discussions related to gender.

Methods

The Python programming language (Van Rossum and Drake 2009) was used to perform this analysis.

Data collection

As mentioned earlier, for UBC salary data, I used the salary PDFs that UBC releases every year (The University of British Columbia 2020). The following steps were taken to collect the data.

- Used the `requests` package to access the UBC Financial Reports [webpage](#).
- Extracted all links and filter for the Statement of Financial Information (SOFI) PDF links which contain salary information.
- If there were any PDFs from which I have not already collected salary data, extracted the text from the PDF using the package `pypdf`.
- Stored the new salary data

An excerpt of the raw salary data is below.

```
Kolhatkar, Ashra 79,036 550 Kolhatkar, Varada 88,579 - Kolind, Shannon H 108,817
```

Data cleaning

In this section, the following steps were taken to clean the salary data:

For each year of salary data,

- Removed special characters from text. For example, §.
- Removed unnecessary text content. For example, the “[Auditor’s] Qualified Opinion”.

- Uses regex to process the raw text data into a structured DataFrame with columns: **Name**, **Remuneration**, **Expenses**. Remuneration will be referred to as “salary” in this report.
- Split the **Name** column into first and last names.
- Converted salary (remuneration) and expenses columns to a numeric data type.
- Shortened first and last names to allow for easier name matching between years. For example, someone’s name in 2020 could be “Bob M Sherbert” and in 2021 their name could be “Bob-M Sherbert”. This name would be shortened to Bob Sherbert to avoid mismatching.

Then, I concatenated dataframes from all years together.

Table 1 shows an expert of the cleaned salary data.

Table 1: Clean UBC Salary Data

	Last_Name	First_Name	Remuneration	Expenses	Year
0	Aamodt	Tor	193153	5597.0	2023
1	Abanto	Arleni	107723	393.0	2023
2	Abbassi	Arash	109136	82.0	2023
3	Abdalkhani	Arman	101829	NaN	2023
4	Abdi	Ali	238203	2981.0	2023

Gender Prediction

Babynome Corpus

In order to predict gender, I used datasets with babynames and their assigned genders. In order to have a somewhat diverse set of baby names, I used babynames from Canadian, American, and Indian sources (Statistics Canada, n.d.; Kaggle 2017; Sharma 2020). I will call this collection of babynome datasets, the “babynome corpus.”

For each UBC staff name, I found whether that name was more common among girls or boys in the babynome corpus. Then, I guessed the gender that was most common. If the name is not present in the corpus, the guessed gender was **None**.

In Table 2, the **Confidence_Score** column shows the percentage of gender majority from the babynome corpus. For example, if 95% of babies named George in the corpus were male, the **Confidence_Score** column value would be 0.95.

The Indian Babynome dataset did not include a count for males and females, so if a name only appeared in the Indian Babynome dataset, an arbitrary **Confidence_Score** of 0.85 was given.

To minimize misgendering, staff names that had less than a 0.8 **Confidence_Score** were given a **Guessed_Gender** of **None**.

Table 2: Babyname Data

	index	Sex_at_birth	First_Name	Confidence_Score
95660	24125.0	Female	Irin	0.5

Table 2 shows the name Irin has a **Confidence_Score** of 0.5. So, since its less than the 0.8 threshold, anyone with that name would have a **Guessed_Gender** value of **None**.

Table 3 shows some of the predictions made on UBC staff members using the babyname corpus.

Table 3: Babyname Data

	First_Name	Guessed_Gender	Confidence_Score
0	Tor	Male	1.0
1	Arleni	Female	1.0
2	Arash	Male	1.0
3	Arman	Male	1.0
5	Yasmine	Female	1.0

Around 91.98% of UBC staff names were matched with names in the babyname corpus. 96.2% of the corpus matches had a **Confidence_Score** over 0.8 and were kept, the rest were given a prediction of **None**.

NLTK

For UBC staff names that were not found in the babyname corpus, I used a natural language processing model to predict genders (Bird, Loper, and Klein 2009).

Table 4 shows some examples of names not found in the babyname corpus.

Table 4: Example Staff Names Not Found in Babyname Corpus

	First_Name
0	Fatawu
1	Tamiza
2	Purang
3	Ninan

Table 4: Example Staff Names Not Found in Babynome Corpus

First_Name	
4	Reto

In order to train and evaluate the model, the babynome corpus was split into a training and test set.

Features used to train the Naive Bayes model were the last 2, 3, and 4 letters of the staff member’s first name.

The accuracy on the test set was 0.86. However, the accuracy on the UBC staff names that were missing from the babynome corpus is very likely lower than 0.86. The data we are making predictions on is quite different from our training data.

Below are the top three features the classifier found most useful for making correct predictions.

Most Informative Features

last_4_letters = 'isha'	Female : Male	=	305.2 : 1.0
last_4_letters = 'etta'	Female : Male	=	193.7 : 1.0
last_3_letters = 'cia'	Female : Male	=	179.7 : 1.0

We can see there are patterns in first names that could be helpful for predicting gender. However, these patterns may not show up often in the unique UBC staff names that were not in the babynome corpus.

Finally, after making predictions on the UBC staff data, we can see in Table 5 the predictions our classifier was least confident about, and in Table 6 and the predictions it was most confident about.

Table 5: Least Confident NLTK predictions

	First_Name	Guessed_Gender	Confidence_Score
1097	Faride	Female	0.43
694	Jiahua	Female	0.44
1646	Xiaohua	Female	0.44
1404	Chuan-Hui	Male	0.44
1092	Qingshi	Female	0.44

Table 6: Most Confident NLTK predictions

	First_Name	Guessed_Gender	Confidence_Score
1501	Sav	Male	0.86
1183	Ninan	Male	0.86
1181	Tamiza	Female	0.86
1179	Bingshuang	Male	0.86
2213	Cicie	Female	0.86

To represent the extra uncertainty in using a classifier compared to the babynome corpus, the **Confidence_Score** column for NLTK predictions is the classifier’s predict-proba score multiplied by 0.86. Where 0.86 is the accuracy of the classifier on the test set.

Like with the babynome corpus predictions, NLTK predictions with a **Confidence_Score** less than 0.8 were given a gender prediction of **None**. This was the case if the predict-proba score from the classifier was less than 0.93.

Overall, 80.2% of the NLTK predictions were over the 0.8 threshold and were kept. The rest were given a prediction of **None**.

Exploratory Data Analysis

Using gender predictions, I created a variety of plots showing the salaries of staff members of UBC. This data only includes staff members making over 75000 CAD annually.




Figure 1: Top Ten Salaries

Figure 1 shows the top ten UBC staff salaries. In 2024, Dawn Jia, female President and CEO of UBC Investment Management Trust, had the highest salary by far. The rest of the top ten salaries were guessed males.




Figure 2: Salary Boxplot

Figure 2 shows box plots of UBC staff salaries split by guessed gender. In 2024, this plot showed the male distribution shifted and skewed towards higher salaries. There were many outliers for males and females.




Figure 3: Salary Histogram

Figure 3 shows a histogram plot of UBC staff salaries. This histogram is split by guessed gender and reflects similar information to the box-plots. In 2024, similar to the box-plots, the male distribution was shifted and skewed towards higher salaries.




Figure 4: Salary Line Plot

Figure 4 shows a line plot of median UBC staff salaries. For both guessed genders, there seems to be fairly minimal change in median salary between 2020 and 2023, and then an increase in median salary in 2024. Males have a higher median salary for (at least) years 2020 to 2024.




Figure 5: Salary Change Line Plot

Figure 5 shows a line plot of the median UBC staff salary percentage change. The data point at 2024 reflects the median percentage change from FY23 to FY24. To calculate these values, first I calculated the salary percentage change for individual staff members who had data for consecutive years. Then, for each year transition, I calculated the median percentage change for guessed males and guessed females.

For both guessed genders, there is a slight decrease in median percentage change between 2020 and 2023, and then a large increase in 2024. This is in line with our observations for Figure 4. Females have a slightly higher median percentage change for years 2020 to 2024.

Limitations

Salary Equality

Figure 4 showed males having a higher median salary from (at least) 2020-2024. Historical wage gaps and gender inequality can make it tempting to draw quick conclusions about salary unfairness from this plot. Even if gender predictions were 100% accurate, it is important not to assume a reason for males having a larger median salary for years 2020–2024 without further investigation. This could be due to a variety of factors, such as work experience or males working in higher-paying fields.

Predicting Gender

Not all the babynames datasets come from reliable sources. Two of them come from Kaggle and it is difficult to ascertain how authentic the data is. It's possible that some of the data is fictitious, which could affect the results of the gender corpus predictions. If I were to redo this project I would spend more time looking for reputable data sources.

Another limitation is that I do not have a good estimate of the accuracy of gender predictions using the NLTK Naive Bayes classification model. For one, the baby name training data is quite different from the UBC data I am predicting on. Additionally, the Naive Bayes model assumes covariate independence conditional on class. The features used are correlated, violating this assumption. This can result in reduced performance and inaccurate class probabilities. Despite the assumption violation, I chose to use the Naive Bayes classifier as it efficiently handles text data, large datasets, and seems to create reasonable predictions. Also, I attempted to minimize the impact of incorrect Naive Bayes classifications by only including those that the model was most confident about.

For simplicity, I chose to only predict male and female genders. However, many people do not fit into this binary and therefore my accuracy for them is 0.

Even if I had the most reliable data sources, the best model, and more gender options, first name is not a direct indicator of gender so misgendering may still occur.

Limitations Result

Overall, due to the limited analysis and unknown accuracy in gender prediction, this report should not be used to draw any conclusions around gender and salaries, especially in regards to salary equality.

Conclusion

Due to the limitations around gender, the most interesting part of this project for me is seeing the percent changes in salary each year. I'm curious as to why there was such a large jump in 2024.

If I were to improve on this project, I would:

- Spend more time researching which classifier would work best. Perhaps there is one that handles text data well and has less violated assumptions.
- Implement cross-validation to get an improved accuracy estimate for the NLTK classifier.
- Do additional research on datasets with east asian names and genders as adding such a dataset to the corpus may help match more names, as well as improve the classifier.

Overall, I really enjoyed learning about data cleaning, classification, and reproducibility.

References

- Bird, Steven, Edward Loper, and Ewan Klein. 2009. *Natural Language Processing with Python*.
- Kaggle. 2017. “US Baby Names.” <https://www.kaggle.com/datasets/kaggle/us-baby-names/data>.
- Sharma, Anany. 2020. “Indian Names Dataset.” <https://www.kaggle.com/datasets/ananysharma/indian-names-dataset>.
- Statistics Canada. n.d. “Table 17-10-0147-01 First Names at Birth by Sex at Birth, Selected Indicators (Number).” <https://doi.org/https://doi.org/10.25318/1710014701-eng>.
- The University of British Columbia. 2020. “UBC Statement of Financial Information.” Open Government - Open Data. <https://finance.ubc.ca/reporting-planning-analysis/reports-and-disclosures>.
- Van Rossum, Guido, and Fred L. Drake. 2009. *Python 3 Reference Manual*. Scotts Valley, CA: CreateSpace.